

CSE 4733/6733 - Operating System 1

S. Torri, T. Ritter

Mississippi State University





Figure 20.1: A 16KB Address Space With 1KB Pages


- Typically have one page table for every process, IE. 100 processes means 100 page tables
- Simple linear table is too memory intensive

Linear Page Tables

1. Can we just make the page bigger, to make the table smaller?
 - Waste within pages - IE. Internal fragmentation
 - End up with MANY fewer jobs in-memory and memory quickly fills up
 - This is a tuneable parameter for a kernel
 - Not that simple...



- Typically have one page table for every process, IE. 100 processes means 100 page tables
- Simple linear table is too memory intensive



- Have a system that supports multiple page sizes
- Linux has *Transparent Huge Pages*[1]




Figure 20.1: A 16KB Address Space With 1KB Pages

- Have a system that supports multiple page sizes
- Linux has *Transparent Huge Pages*[1]

2023-02-21

Linear Page Tables

1. Complications?

- A single TLB entry will map to a much larger amount of virtual memory in turn reducing the TLB misses.
- Now use a *khugepaged* daemon to scan memory and collapse sequences of std pages into huge pages
- Might be worth the effort for a database or other special purpose system, that why we have the libraries in Linux[2]
- Virtual memory management becomes more complicated, latency spikes as OS moves data around and updates page entries...

- Create a hybrid page table with both segmentation and paging
- User address space broken into a number of segments, with each segment into pages
- Both have strengths, paging to remove external fragmentation, segmentation lends itself to protection and sharing

└ Hybrid Page Table

1. Comments

- Easy base address and length calculation to protect memory bounds
- Combined segmentation and paging has advantages, but brings back external fragmentation

- ❖ Create a hybrid page table with both segmentation and paging
- ❖ User address space broken into a number of segments, with each segment into pages
- ❖ Both have strengths, paging to remove external fragmentation, segmentation lends itself to protection and sharing


Advanced Page Tables

CSE 4733/6733


2023-02-21

Advanced Page Tables


Advanced Page Tables



Multi-level Page Tables



Multi-level Page Tables



1. Note: Classic Trade-Off

- We created a OS data structure with a smaller memory footprint
- We see that a TLB *MISS* is more expensive than before
- More page table levels are possible, what if the Page Directory is too large for a single page???
- Supports sparse address spaces well, and page-table space is proportional to address space
- Subtle note, valid flag in a Page Pirectory mean there is at least one page entry valid in the page where it points, unlike in the page table itself
- More complexity in the design...

Inverted Page Table

- Create a *single* OS page table that has one entry for each real memory frame (physical page)
- Entry contains PID and what virtual page maps to that frame
- Create a simple hash function to speed up the linear search
- Inverted since the table is indexed by real memory frames, not virtual
- Page tables are a data structure that allows for trade-offs

└ Inverted Page Table

1. Requirements

- Table requires PID of the process with the frame
- Table requires control bits for valid, modified (dirty), protection, etc.
- Table requires a chain pointer since your hash may point to the same entry as another virtual address

- Create a *single* OS page table that has one entry for each real memory frame (physical page)
- Entry contains PID and what virtual page maps to that frame
- Create a simple hash function to speed up the linear search
- Inverted since the table is indexed by real memory frames, not virtual
- Page tables are a data structure that allows for trade-offs

└ Swapping

1. Comments

- Page faults are expected, really more of a page miss
- Hardware will raise an exception, and the OS will have to swap the page requested into memory
- The OS will have a page-fault handler to service page faults

- Multiprogramming, large jobs all drive an OS to support swapping
- Swap space is secondary storage, typically a disk address/chunk of space
- Need another flag - a *present bit* to determine if a page is in memory or swap
- The terminology of swapping - A *page fault* is when the OS wants to access a page that is not in memory

- Swapping is by nature a *slow* process compared to memory access times, the TLB is a memory cache
- This is not a hardware function, rather the OS is going to do it
- When memory is full, you have to have a Page Replacement Policy

└ Swapping

1. Comments

- Page faults are expected, really more of a page miss
- Hardware will raise an exception, and the OS will have to swap the page request into memory
- The OS will have a page-fault handler to service the page fault

- ❖ Swapping is by nature a slow process compared to memory access times, the TLB is a memory cache
- ❖ This is not a hardware function, rather the OS is going to do it
- ❖ When memory is full, you have to have a Page Replacement Policy

- Kernel parameters of: *high watermark (HW) and low watermark (LW)*
- Linux has a *kswapd* to swap pages as needed
- When low watermark is reached then kswapd wakes to reclaim memory in the background.
- It stays awake until free memory reaches the high watermark.

2023-02-21

└ Swapping

1. Comments

- Linux kernel tuning is a complicated enterprise [3] eg. `sysctl -a`
- The article quoted 49 Virtual Memory parameters out of 1,200 you can tune
- Linux has a 3-level page table memory management scheme and is complex
- Swapping is normal, but low memory can make a machine work at the speed of disk I/O rather than memory I/O

└ References

2023-02-21

References I



[Free Software Foundation.](#)

Transparent hugepage support, 2023.



[Erik Rigtorp.](#)

Using huge pages on linux, 2020.



[Al Williams.](#)

Fine tuning linux virtual memory, 2022.

Free Software Foundation.
Transparent hugepage support, 2023.
Erik Rigtorp.
Using huge pages on linux, 2020.
Al Williams.
Fine tuning linux virtual memory, 2022.